

Course wrap-up · AI-Augmented Productivity for Finance

Two days, five sessions, one tool on your GitHub — and the habits that make an AI output trustworthy.

The storyline of the course

- | | |
|------------------------------|--|
| 1 · grounded prompts | the model answers only from evidence you supply |
| 2 · valuation vs peers | Apple valued against 7 peers, live from EDGAR |
| 3 · debugging & verification | the valuation repaired; every quote verified |
| 4 · the screening workflow | 16 companies in, a justified shortlist out |
| 5 · the governed agent | the model plans its steps, inside limits you wrote |

Each session produced a working artifact; together they are one tool.

The division of labor that carried every session

- **Code does the math** — pandas and numpy are deterministic; the model is not
- **The model does judgment, in words** — grounded in the data you supply, nothing else
- **Code checks the words** — schemas, sanity tests, evidence verification, numeric audits
- **A human signs** — the gate before anything is saved, sent or published

What makes an AI output trustworthy

- **Grounding:** the model saw only your data, so every claim traces to a source
- **Validation at the boundary:** typed contracts — malformed output cannot pass
- **Measured quality:** a forecast ships with its holdout error; a quote with its evidence
- **Limits in code:** steps, budgets, tool whitelists, audit trails — autonomy granted like a mandate

The toolbox you met

Tool	What it does	Sessions
Claude API	drafts, explains, reasons — inside your rules	all
pandas · numpy	deterministic math: screens, valuations, forecasts	2–4
SEC EDGAR	primary-source fundamentals, free	2–5
Pydantic	typed contracts at the model boundary	3–5
LangGraph	workflows and agents as explicit graphs	4–5
PydanticAI	the governed agent loop, packaged	5
Skills (SKILL.md)	procedures the agent loads on demand	5
git · GitHub	your work, versioned and public	2–5

Your capstone submission

- **What:** your repository URL + a one-page memo inside the repository
- **The memo:** the job · a real run on your own pair · the trust story · one honest number · the next step
- **Grading:** the trust story carries the most weight — a documented, caught failure is stronger evidence than an unexamined success
- **How:** post your repository URL in the class chat at the end of this session

Brief and rubric: `session-05-agents/capstone/capstone-brief.md`